

Greenscape Pro — AI Agent Strategy

Prepared for License & Scale · Take-home submission · Abdullah Omar

Marcus runs a \$4.2M hardscape business with one large revenue hole and several smaller ones. The biggest dollar gap is the proposal cycle. The biggest *strategic* opportunity is different: an agent that can ship in two weeks, recover roughly \$784K in latent revenue, and earn the trust we'll need for the structural changes that come after. The five agents below are ordered by what gets built first, not by what's largest.

1. Closed-Lost Re-engagement Agent

Purpose. Mine 1,400 dormant leads with personalized outreach that sounds like Marcus, not a blast.

- Pulls each closed-lost lead's GHL notes, project type, last touch date, and historical objection
- Generates a short personalized email or SMS that references what the customer originally wanted
- Routes every draft to Marcus for one-tap approval; nothing sends without his sign-off
- Sends through GHL on a staggered schedule, no more than one touch per lead per quarter
- Routes any reply back to Brittany for human follow-up before quoting begins

Replaces. Brittany's ad-hoc re-engagement attempts, which Marcus described as "not systematic."

ROI. $1,400 \text{ leads} \times 2\% \text{ reactivation} \times \$28\text{K average} = \sim\$784\text{K of latent revenue}$. CAC is already paid. The leads are already cold. The downside is essentially zero.

Why #1. Bigger absolute dollars exist on this list — the Quote Agent is roughly 3× larger. But #1 is a sequencing question, not a size question. Closed-lost ships in two weeks. It carries no risk of disrupting active sales process. It generates the response data and trust capital that make the harder structural agents (especially #2) easier to fund and deploy. Senior consultants don't lead with the biggest swing. They lead with the cleanest defensible win. This is that.

2. Quote Acceleration Agent

Purpose. Ship a customer-ready proposal within 24 hours of the site walk, with Marcus as reviewer instead of author.

- Ingests Marcus's site-walk notes (text or voice memo) plus customer context from GHL
- Extracts structured scope, matches against the 200+ line-item pricing catalog, flags low-confidence items
- Drafts the proposal in Marcus's voice; auto-triggers a Carlos render request when total exceeds \$30K
- Routes to Marcus for one-click review and edit; nothing sends without approval
- On approval, generates the PDF and pushes it to the customer through GHL

Replaces. Marcus's 2–3 hours of personal proposal-drafting per deal, plus the 6–9 day cycle that follows.

ROI. 35–40% of qualified leads currently go to faster competitors. At ~150 won deals/yr × \$28K × ~37% lost share, this is **~\$2.2M–\$2.8M/yr in addressable revenue** — the largest single dollar gap in the business.

Why #2 and not #1. This is the biggest opportunity. It's also the structural change that asks Marcus to stop authoring proposals — the part of the business he most identifies with. That conversation goes much better after #1 has shipped and proven the model. Build the trust first. Then ask for the keys to the kingdom.

3. Post-Sign Acceleration Agent

Purpose. Stop the 4–6 week stall between contract signed and crew on site.

- Tracks each project through HOA, permit, and deposit stages with stage-aware nudges
- Sends reminders to customer, HOA, and city through GHL on the right cadence
- Escalates to Jenna when a milestone slips beyond its expected window
- Reports project status to Marcus weekly so he can plan crew schedules with real data

Replaces. Jenna's manual chase loop on 8–12 active projects.

ROI. 8–12 projects × \$28K = **\$224K–\$336K of cash perpetually delayed.** Compounds because every week of HOA limbo is a week the crew can't be on the next job.

Why here. Operationally significant, but it accelerates revenue Greenscape has already booked rather than creating new revenue. Worth doing, ranked below dollars they don't book at all.

4. Build-Phase Communication Agent

Purpose. Make every customer feel like Marcus is personally updating them, even on the 70% of jobs where he isn't.

- Triggers off CompanyCam photo uploads and Jobber milestone events
- Drafts a short warm update in Marcus's voice for the customer
- Marcus approves with one tap; sends through GHL
- Logs every send so anxious customers don't get five identical messages on the same day

Replaces. Marcus's inconsistent personal Loom updates (~30% of jobs today) and the inbound "what's going on?" anxiety calls Jenna fields.

ROI. Harder to quantify cleanly, but Marcus told us referrals come from this exact behavior. If 5 extra referrals per year close at \$28K, that's \$140K — and it eliminates the support load on Jenna in the process.

Why here. Compounds back into the lead engine through referrals, which makes it a multiplier on every other agent below it.

5. Approvals Co-Pilot

Purpose. Be the rule book Jenna asked for out loud on the call.

- Sits in Slack, answers the 5–10 daily approval questions Jenna currently sends Marcus
- Pulls from a documented framework Marcus confirms once (change-order pricing, refund policy, scope-add rules)
- Self-serves the 80% that follow the rules; escalates the rest to Marcus with full context
- Logs every decision so the framework gets sharper over time

Replaces. The 5–10 Slack pings/day from Jenna to Marcus, plus the wait time those pings impose on her.

ROI. ~150–300 hours of Marcus's executive time per year, directly hitting his closing line: *"I'm tired. I want my evenings back."*

Why #5. Real value, but it's time savings, not revenue. It's also the riskiest agent here — only as good as the rules Marcus gives it, and he hasn't written them down yet by his own admission. Worth shipping once the bigger leaks are sealed.

The binding constraint nobody flagged on the call

Every agent on this list has a human-in-the-loop step where Marcus approves before send. Ship five agents and we've made Marcus babysit five things — the exact thing he told us he doesn't want: *"I do not want a bunch of fancy tech I have to babysit."*

The fix isn't to drop human-in-the-loop, because Marcus needs control. The fix is portfolio architecture. Every agent routes outputs to a single approval surface, ranked by urgency and dollar value, with confidence thresholds calibrated so high-trust drafts auto-send and only edge cases bubble up. That converts "five agents to babysit" into ten minutes a day of executive review. This isn't a feature of any single agent. It's the design rule the whole portfolio has to obey, and it should be locked in before agent #2 ships.

Why my #1 isn't the founder's #1

Marcus's stated #1 is "speed up quoting." That's my #2. I'm leading with closed-lost re-engagement for sequencing reasons rather than dollar reasons.

His framing of the quote agent — "speed up quoting" — is also too small even when we get to it. The agent doesn't make Marcus faster. It removes him from the critical path entirely. Marcus is the bottleneck, and you don't unblock a bottleneck by making it faster. You route around it. The unlock is "Marcus reviews, the agent authors." That's a structural change, and structural changes go better after Marcus has watched a smaller agent earn its keep.

Two of his other stated priorities don't survive the math:

- **Crew coaching (his #3)** is \$104K/yr against \$2.8M/yr in lost quote revenue. Order of magnitude smaller. It bugs him because he watches the small mistakes happen, but emotion isn't the priority signal. Cut.
- **Marketing/content (his #4)** solves a constraint that doesn't bind. Marcus disqualified it on the call himself: *"Are you lead-volume constrained or quote-volume constrained?" "Quote."* Building this right now optimizes the wrong variable.

One agent I considered and cut

Daily Crew Check-In Summarizer. It was on Marcus's "fire myself from" list and I see the appeal — let him skip reading four daily check-ins. But it's pure time savings: no new revenue, no cash acceleration, no referrals. Ranking it above any of the five above would optimize the wrong variable. Phase 2.

Interdependencies worth flagging

- The Build-Phase Comms agent (#4) feeds the closed-lost engine (#1) through referrals — happy mid-build customers refer the next batch of leads that go cold and re-enter the funnel 18 months later.
- The Approvals Co-Pilot (#5) reclaims Marcus's review attention, which is the binding constraint on every other agent's approval loop.
- The Quote Agent (#2) only becomes politically viable after #1 ships. Marcus has to watch an agent draft customer-facing content well before he'll hand over proposal authoring.

Assumptions I made

- The 35–40% lost-to-speed share refers to qualified leads (post-site-walk), per the auditor's call notes
- "Marcus's voice" can be approximated with a system prompt and a handful of his actual messages as examples; no fine-tune needed
- GHL has a usable API for SMS/email send and reading lead notes; if it doesn't, we wrap it with a manual approval queue and revisit
- The 2% closed-lost reactivation rate is conservative given the personalization quality the agent can deliver; I'd model upside at 3–4%